Abstract

. . .

Contents

1   Introduction

2   Contributions

We study the Dynamic Balanced Graph Partitioning problem and show a competitive ratio of $O(k \log k)$ for our component-based deterministic algorithm CREP in the variant ADJ-DEL which runs in polynomial time. We also discuss challenges in the competitive analysis of this algorithm and of another algorithm, CORE-DEL, for which we have not shown a competitive ratio. Both approaches are adaptations of prior work ([2, 5]).

   We discuss the details of our implementation which uses a decomposition tree data structure in order to reduce the update times when recomputing the components upon receiving new requests.

   We evaluate our implementation and compare it with CORE-DEL that is also based on components, but for which we have not shown a competitive ratio. Furthermore we compare it with two existing implementations in the METIS and ParMETIS framework and we explore first ideas of adjusting parameters to improve the results of the presented algorithms even further.

3   Definitions and Notation

In this section we introduce several definitions and notations that will be used throughout the thesis.

   First we define a graph $G = (V, E, [w])$ where V is a set of vertices and $E \subseteq V \times V$ is a set of edges. $w : E \to \mathbb{N}$ assigns each edge an (integer) weight. Note that we always assume that a graph is undirected, i.e. edges are of the form $\{u, v\}$ where $u \neq v$ and $u, v \in V$. A graph is only unweighted or directed if it is explicitly mentioned or clear from the context.

   Given a graph $G = (V, E, w)$ we define an (edge) cut of G as a pair of two disjoint subsets X, Y of V such that $X \cup Y = V$. The value of this cut is then the sum of the weight of edges between nodes from X and Y, i.e. $\sum_{e=\{u,v\} \in E : u \in X, v \in Y} w(e)$ is the value of the cut (X, Y). Note that such a cut can also be defined by the set of the edges connecting X and Y that are cut. We call a cut a *minimum (edge) cut* of G if it is one of the cuts with minimum value.

   We use this definition a minimum cut to introduce the notion of *connectivity*. The connectivity of a graph G is equal to the value of a minimum edge cut of G. This definition will be used in order to define the communication components our algorithm maintains as these are subsets of V who induce subgraphs of high connectivity. We explain the concept of components in greater detail in Section 6.

   Furthermore we define the term (s, t)-*cut* a cut (X, Y) for which $s \in X$ and $Y \in X$, i.e. a (s, t)-cut separates the nodes s and t in G. Then a *minimum (s, t)-cut* is a (s, t)-cut of minimum value. Note that a minimum (s, t)-cut is not necessarily a minimum cut.

4 Problem Definition

Avin, Bienkowski, Loukas, Pacut and Schmid ([5]) first introduced the Dynamic Balanced Graph Partitioning problem under the name of *Balanced RePartitioning* (BRP) problem, but changed the name to the former in a more recent version of their paper ([2]).

The task is to maintain a partitioning of a dynamic graph consisting of $n = k \cdot l$ nodes that communicate with each other into l parts, each of size k while minimizing both the cost due to communication and due to node migrations defined as follows. The communication cost is zero if both nodes are located on the same server at the time the request needs to be served and it is normalized to one if they are mapped to different servers. An algorithm may perform node migrations in order to change the mapping of nodes to servers prior to serving the communication request at time t. Such a move of one vertex incurs cost $\alpha > 1$.

More formally we are given l servers $V_0, ..., V_{l-1}$, each with capacity k and an initial perfect mapping of $n = k \cdot l$ nodes to the l servers, i.e. each server is assigned exactly k nodes. An input sequence $\sigma = (u_1, v_1), (u_2, v_2), ...(u_i, v_i), ...$ describes the sequence of communication requests: the pair (u_t, v_t) represents a communication request between the nodes u_t and v_t arriving at time t. At time t the algorithm is allowed to perform node migrations at a cost of $\alpha > 1$ per move. After the migration step the algorithm pays cost 1 if u_t and v_t are mapped to different servers and does not pay any cost otherwise. Note that an algorithm may also choose to perform no migrations at all.

We are in the realm of competitive analysis and as a result we compare an online algorithm ONL to the optimal offline algorithm OPT. ONL only learns of the requests in the input sequence σ as they happen and as a result only knows about the partial sequence $(u_1, v_1), ..., (u_t, v_t)$ at time t whereas OPT has perfect knowledge of the complete sequence σ at all times.

The goal is to design an online algorithm ONL with a good competitive ratio with regard to OPT defined as follows.

An online algorithm ONL is $\rho - competitive$ if there exists a constant β such that

$$\text{ONL}(\sigma) \leq \rho \cdot \text{OPT}(\sigma) + \beta \; \forall \sigma$$

where $\text{ONL}(\sigma)$ and $\text{OPT}(\sigma)$ denote the cost of serving input sequence σ of ONL and OPT respectively.

Often we allow the online algorithm to use larger capacities per server. In this case we speak of an *augmentation* of δ in the case where the online algorithm is allowed to assign $\delta \times n/k$ nodes to each server where $\delta > 1$. This augmented online algorithm is than compared with the optimal offline algorithm OPT which is not allowed to use any augmentation.

5 Related Work

With the ever increasing importance of distributed computation systems also come new challenges. One such challenge is an increase in network traffic. This increased network load may impact the performance of cloud computing networks significantly

and should hence be minimized. For example studies ([9, 10]) have shown that reconfigurable datacenter networks can provide performance similar to regular datacenter networks while having significantly lower cost in some cases.

One avenue of reducing this network load and hence also increasing the performance of cloud applications is to relocate the different communication partners inside the network dynamically in order to reduce the volume of inter server communication.

5.1 Dynamic Balanced RePartitioning

Avin et al.[2] initiated the study of the online variant of the Balanced RePartitioning (BRP) problem that is the topic of this thesis. They propose a deterministic algorithm for the Dynamic Balanced Graph Partitioning problem with augmentation $2 + \epsilon$ for any $\epsilon > 1/k$. They also show a lower bound of $k - 1$ for the competitive ratio of any online algorithm for the Dynamic Balanced Graph Partitioning problem on two clusters via a reduction to online paging. Furthermore they show that no δ-augmented deterministic online algorithm can achieve a competitive ratio smaller than k for any augmentation $\delta < l$.

5.2 Restricted Variants of Balanced RePartitioning

Restricted variants of the Balanced RePartitioning problem have also been studied. Here one assumes certain restrictions of the input sequence σ and then studies online algorithms for these cases.

Avin, Cohen, Parham and Schmid ([3]) study one such case: the authors assume that an adversary provides requests according to a fixed distribution of which the optimal algorithm OPT has knowledge while an online algorithm that is compared with OPT has not. Further they restrict the communication pattern to form a ring-like pattern, i.e. for the case of n nodes $0, ..., n-1$ only requests r of the form $r = \{i \mod n, (i+1) \mod n\}$ are allowed. For this case they present a competitive online algorithm which achieves a competitive ratio of $O(\log n)$ with high probability.

Henzinger, Neumann and Schmid ([12]) studies a special *learning variant* of the Dynamic Balanced Graph Partitioning problem specified above. In this version it is assumed that the input sequence σ eventually reveals a perfect balanced partitioning of the n nodes into l parts of size k such that the edge cut is zero. In this case the communication patterns reveal connected components of the communication graph of which each forms one of the partitions. Algorithms are tasked to *learn* this partition and to eventually collocate nodes according to the partition while minimizing communication and migration costs.

The authors of [12] present an algorithm for the case where the number of servers is $l = 2$ that achieves a competitive ratio of $O((\log n)/\epsilon)$ with augmentation ϵ, i.e. each server has capacity $(1 + \epsilon)n/2$ for $\epsilon \in (0, 1)$.

For the general case of l servers of capacity $(1 + \epsilon)n/l$ the authors construct an exponential-time algorithm that achieves a competitive ratio of $O((l \log n \log l)/\epsilon)$ for $\epsilon \in (0, 1/2)$ and also provide a distributed version. Additionally the authors describe a polynomial-time $O((l^2 \log n \log l)/\epsilon^2)$-competitive algorithm for the case

with general l, servers of capacity $(1 + \epsilon)n/l$ and $\epsilon \in (0, 1/2)$.

It is important to stress that the assumption that the requests reveal a perfect partitioning of the communication nodes is not applicable for most practical applications and thus it is important to study the general BRP problem without restricting σ.

5.3 Clustering

Clustering is the process of generating subsets of elements with high similarity ([11]). Desirable properties are *homogeneity*, i.e. high similarity within elements of a cluster and that the similarity is low between elements of different clusters, a property called *separation*. Hartuv and Shamir ([11]) approach this problem by constructing a similarity graph and then separating the nodes via repeated computations of minimum edge cuts. This approach is quite similar to the one we use in order to determine sub-graphs of high connectivity of the communication graph induced by the requests.

However we consider an online problem, i.e. algorithms need to react dynamically to changes in the graph and need to maintain their data structures and adapt accordingly whereas clustering considers complete data sets which are static.

5.4 Online Paging

In the Online Paging problem ([8], [7]) one is given a scenario with a fast cache of k pages and $n - k$ pages in slow memory. In an online manner pages are requested, i.e. without prior knowledge of future requests. If a requested page is in the cache at the time of the request it can be served without cost. If it is in slow memory however, then a *page fault* occurs and the requested page needs to be moved into the cache. If the cache is full then a page from the cache needs to be evicted, i.e. moved to the slow memory in order to make space for the requested one. The goal is to design algorithms which minimize the number of such page faults.

The Dynamic Balanced Graph Partitioning problem can be seen as a generalization of Online Paging. In fact, Avin et al. ([2]) have shown a reduction of Online Paging to this problem.

However the standard version of Online Paging has no equivalent to the option of serving a request remotely as is possible in the Dynamic Balanced Graph Partitioning problem. The variant *with bypassing* allows an algorithm to access pages in slow memory without moving them into the cache, thus providing such an equivalent. It is worth stressing however that in our problem requests involve two nodes while in Online Paging the nodes themselves are requested.

5.5 Static Balanced Graph Partitioning

The Static Balanced Graph Partitioning problem is the static offline variant of the topic of this thesis. In this version an algorithm may not perform any migrations, but has perfect knowledge of the request sequence σ and then needs to provide a perfectly balanced partitioning of the $n = k \cdot l$ nodes into l sets of equal size k.

This scenario can be modelled as a graph partitioning problem where the weight of an edge corresponds to the number of requests between its end points in the input sequence σ. An algorithm then has to provide a partition of the nodes into sets of exactly k nodes each while minimizing the total edge weights between partitions, i.e. an algorithm needs to minimize the edge cut of the graph.

For the case where $l \geq 3$, Andreev and Räcke have shown that there is no polynomial time approximation algorithm which guarantees a finite approximation factor unless P=NP ([1]).

6 Algorithmic Ideas

In this section we describe two different solution approaches to the Dynamic Balanced Graph Partitioning problem. We call these approaches CORE-DEL and ADJ-DEL respectively. We first describe this general approach that is common to both algorithms and then address the specific differences and analysis ideas.

Both methods share a similar concept at their core: a second-order partitioning of the communication nodes into *communication components* which represent node-induced sub-graphs of the original communication graph given by the requests from the input sequence σ. As more requests from σ are revealed to the algorithms they merge the corresponding components once they are suitably connected and relocate the nodes of the new component in such a way that all the nodes of a component are always located on the same server.

More formally, initially each node forms a singleton component, but as the input sequence σ is revealed to the algorithms new communication patterns unfold. The algorithm keeps track of these patterns by maintaining a graph in which the nodes represent the actual communication nodes and the weighted edges represent the number of communication requests between nodes that were part of different components at the time of the request, i.e. for edge $e = \{u, v\}$, $w(e)$ represents the number of paid communication requests between u and v. We say that a communication request between nodes u and v is *paid* if the nodes are located on different servers at the time of the request.

Both algorithms merge a set S of components into a new component C if the connectivity of the component graph induced by the components in S is at least α. After each edge insertion the algorithm checks whether there exists a new component set S with $|S| > 1$ which fulfils this requirement.

If after any request and the insertion of the resulting edge the algorithm discovers a new subset S of nodes whose induced subgraph has connectivity at least α and which is of cardinality at most k it merges the components that form this set into one new component and collocates all the nodes in the resulting set on a single server. If the resulting component has size at least $2/\epsilon$ the algorithm reserves additional space $\min\{\lfloor \epsilon \cdot |C| \rfloor, k - |C|\}$, otherwise the reservation is zero. This reservation guarantees that nodes are not migrated too often for the analysis to work. Note that this also limits the total space a component can use to a maximum of k. This makes sense as a component whose size exceeds k is never merged but instead deleted and hence there would be no benefit to a component taking space more than k.

The collocation of such component sets of at most k individual communication

Algorithm 1 DynamicDecomp

  Initialize an empty graph on n nodes
  turn each of the n nodes into a singleton component
  for all $r = \{u, v\} \in \sigma$ **do**
    if $comp(v) \neq comp(u)$ **then**
      $w(\{u, v\}) \leftarrow w(\{u, v\}) + 1$
    end if
    if $\exists$ component set X with connectivity at least α and $|X| > 1$ and $nodes(X) \leq k$ **then**
      merge(X) and update reservations
    end if
    if $\exists$ component set Y with connectivity at least α and $nodes(Y) > k$ **then**
      delete(Y)
    end if
  end for

nodes is always possible without moving a node not in C due to the allowed augmentation of $2 + \epsilon$. This guarantees by an averaging argument that there is almost at least one cluster with capacity at least k which a newly merged component can be moved to.

If the subset has cardinality greater than k the resulting component is deleted. The definition of this deletion process is the main differentiating factor between our algorithms which we discuss in the following subsections. The common part of both algorithms is also summarized in the form of pseudocode in Algorithm 1. Please note that the subroutine *delete(Y)* of a component set Y is different for each of the algorithms. We also describe the particular challenges each approach entails when it comes to the competitive analysis.

These approaches are fairly similar to the algorithms defined in previous work ([2], [5]). The main differentiating factor is that we merge once a component set reaches connectivity α while the approaches by the other authors do so once the component set reaches a certain density threshold. More specifically they merge a component set S once it fulfills $w(S) \geq (|S| - 1) \cdot \alpha$ where $w(S)$ denotes the cumulative weight of the edges between nodes contained in the components of S. The similarities are especially apparent when comparing the respective lemmas and properties that are used in order to bound the weight between partitions of mergeable component sets. These are lemma 4.3 in [2], property 3 in [5] and Lemma 6 in this thesis.

Now we describe the differences in the deletion steps of CORE-DEL and ADJ-DEL and present the implications for their respective competitive analysis.

6.1   CORE-DEL

We address CORE-DEL first. In this version edges are reset which are *contained* in the deleted component, i.e. all edges $e = \{u, v\}$ are reset to zero if both u and v were contained in component C at the time of its deletion. This approach resembles the one suggested by Avin et al. ([2]) and is also written in pseudocode in Algorithm 2. The idea of the analysis is then to relate the cost of CORE-DEL with the cost of

Algorithm 2 delete(Y) of CORE-DEL
___
   for all $e = \{u, v\} \in E$ **do**
     if $u \in Y$ and $v \in Y$ **then**
       $w(e) \leftarrow 0$
     end if
   end for
___

OPT by considering the respective costs due to requests from a deleted component C in the solution of CORE-DEL as these are of high connectivity and are impossible for OPT to collocate on one server as each deleted component contains more than k nodes. Then one could sum these costs over all such deleted components in order to bound the costs. For this approach to work however one would have to find a way to cleanly separate requests belonging to one deleted component from those belonging to another in order to establish a lower bound on the cost of OPT.

In this case the edges which are adjacent to a deleted component C remain even after the deletion of C and may then contribute the the creation of a new component D. It is now very challenging to attribute any significant cost to OPT for these requests.

Figure 1 shows an example sequence of requests for which this approach does not work. The diagram shows horizontal lines, each representing one of the vertices. A vertical line represents a communication request between its end points. For example the sequence shown contains a communication request between nodes 1 and 2 at time $t = 5$. Now consider the case where $\alpha = 3$ and $k = 3$.

In this sequence the first two requests between nodes 0 and 4 happen without leading to a merge. The following 12 requests lead to a merge of a new component C consisting of the nodes 1,2,3 and 4 which gets deleted by CORE-DEL at time $t = 14$. Note now that the edges corresponding to the requests between nodes 0 and 4 are still present even after this deletion. Finally the request at time $t = 15$ leads to a merge of nodes 0 and 4.

One can see that such cases may also happen at a much larger scale, where almost all requests happen much earlier than the time at which an actual merge of the nodes involved in the requests happens.

6.2 ADJ-DEL

Algorithm 3 delete(Y) of ADJ-DEL
___
   for all $e = \{u, v\} \in E$ **do**
     if $u \in Y$ or $v \in Y$ **then**
       $w(e) \leftarrow 0$
     end if
   end for
___

The second algorithm, ADJ-DEL, resets all the edges contained in the deleted component C but also resets the weights of edges *adjacent* to C, i.e. all edges $e = \{u, v\}$ are reset to zero if u or v were contained in component C at the time of

Figure 1: illustration of the analysis problem with the approach CORE-DEL

Figure 2: illustration for the ADJ-DEL approach

its deletion. This second version is very similar to the algorithm proposed by Avin et al. in [5]. The deletion method is also described in pseudocode in Algorithm 3.

The idea for the analysis is once again to relate the cost of both OPT and ADJ-DEL to the deleted components in the solution of ADJ-DEL.

The fact that ADJ-DEL also resets adjacent edges means that we can uniquely identify requests with the deleted component whose deletion led to the reset of the corresponding edge weights to zero. This means that we do not have to face the problems we discussed in the section on CORE-DEL. However the question remains whether the deletion of adjacent edges is valid, i.e. whether we can still preserve a good competitive ratio for this approach.

Figure 2 illustrates an example of this challenge mentioned above for the ADJ-DEL approach.. In the illustration it is assumed that $\alpha = 3$ and $k = 3$. In this case the input sequence leads ADJ-DEL to first merge nodes 1 and 2 at time $t = 3$ and then

to add node 3 to the resulting component at time $t = 6$. The next two requests do not quite lead to the merge of the node 4 with the component. Instead a series of requests follows where node 1 communicates with other nodes that are outside of its component without any merges. Note that this sequence can be extended until node 1 has communicated with every node except from the nodes $1, 2, 3, 4$. Finally the first 4 nodes are merged at time $t = 15$ at which point the resulting component is deleted as well as *all* edges in the sequence. This means that the amount of adjacent edges that are reset may greatly exceed the amount of edges inside a deleted component. We show in our analysis in Section 7 that this approach is valid and allows for a competitive ratio of $O(k \log k)$.

7   Analysis

For the sake of this analysis we examine the ADJ-DEL algorithm from Section 6 and name the algorithm CREP from now on.

We analyse the competitive ratio of CREP with augmentation $(2 + \epsilon)$ and show in Theorem 14 that CREP is $O(k \log k)$-competitive.

7.1   Algorithm Definitions

We begin our analysis by introducing two general definitions that we will use throughout the analysis.

Definition 1. Define for any subset S of components $w(S)$ as the total weight of all edges between nodes of S.

Definition 2. Let a set of components of size at least 2 and of connectivity α be a *mergeable* component set.

7.2   Structural Properties

Note: These properties are changed to use the connectivity-based approach which generally simplifies them, but guarantees slightly less minimum edge weight within mergeable component sets.

Definition 3. An α-*connected component* is a maximal set of vertices that is α-connected.

The following lemma states that there is always at most one mergeable component set after the insertion of a new edge which CREP then merges. The earliest point in time a new mergeable component set can emerge is after the next edge is inserted. This lemma is very similar to lemma 4.1 from [2] in essence. The only difference lies in the definition of a mergeable component set which we have defined via a connectivity measure while the authors of [2] used density.

Lemma 4. At any time t after CREP performed its merge and delete actions all subsets S of components with $|S| > 1$ have connectivity less than α, i.e. there exist no mergeable component sets after CREP performed its merges.

9

Proof. We proof the lemma by an induction on steps. The lemma holds trivially at time 0.

Now assume that at some time $t > 0$ the lemma does not hold, i.e. there is a subset S of components with connectivity at least α and $|S| > 1$. We may assume that t is the earliest time for which S has connectivity α.

Then the incrementation of the weight of edge e at time t raised the connectivity of S, but S was not merged into a new α-connected component C. if no new component was created at time t then we arrive at a contradiction as CREP always merges if there exists a mergeable component set.

Now assume that a component C was created at time t. This means that C must also contain the endpoints of e. But then the conjunction of C and S forms an even larger subset of components with connectivity at least α which is a contradiction to the maximality of C and S. □

The following lemma is adapted for our connectivity-based approach from corollary 4.2 in [2].

Lemma 5. Fix any time t and consider weights right after they were updated by CREP but before any merge or delete actions. Then all subsets S of components with $|S| > 1$ have connectivity at most α and a mergeable component set S has connectivity exactly α.

Proof. This lemma follows directly from Lemma 4 as connectivities can only increase by at most 1 at each time t and Lemma 4 guarantees that mergeable component sets are merged by CREP directly after they emerge, before a new request is revealed. □

The following two lemmas combined give us a result similar to lemma 4.3 in [2]: bounds on the edge weight that is cut when partitioning a mergeable component set, i.e. a set of components of connectivity at least α.

We start by establishing a lower bound on this edge weight in the following lemma.

Lemma 6. Given a mergeable set of components S and a partition of S into $g > 1$ parts $S_1, ..., S_g$. Then the weight between the parts of the partition is at least $g/2 \cdot alpha$.

Proof. We construct a graph G with the different parts $S_i, i \in \{1, ..., g\}$ of the partition as nodes. Note that this graph is α-connected. We insert an edge for each edge between the parts S_i. Now consider the sum of the weighted degrees of all such nodes S_i in the constructed graph:

$$\sum_{i \in \{1,...,g\}} deg_G(S_i) = 2 \sum_{e \in G} w(e)$$

The equality follows as the left sum counts each edge twice, once for each endpoint. Now consider the fact that each node S_i must have degree at least α with respect to the edges in G because G is α-connected. Hence

$$2 \sum_{e \in G} w(e) = \sum_{i \in \{1,...,g\}} deg_G(S_i) \leq \sum_{i \in \{1,...,g\}} \alpha = g \cdot \alpha$$

10

which gives us that $\sum_{e \in G} w(e) \geq g/2 \cdot \alpha$. $\qquad\square$

In the following lemma we establish the upper bound on the cut edge weight when partitioning a mergeable set of components S into $g \geq 2$ parts.

Lemma 7. Given a mergeable set of components S and a partitioning of S into $g \geq 2$ parts $S_1, ..., S_g$. The weight between the parts S_i is at most $(g-1) \cdot \alpha$ during the execution of CREP.

Proof. Similarly to before e construct a graph $G = (V, E)$ with the different parts $S_i, i \in \{1, ..., g\}$ of the partition as nodes and we insert an edge for each edge between the parts S_i. Note again that this graph is α-connected. We iteratively partition G into subsets via minimum cuts with regard to edge weight, i.e. we consider a minimum edge cut of G which partitions the nodes of G into the subsets V_1 and V_2 and iteratively partition the resulting sets until all sets contain only one node of G each. As this required at most $|V| - 1$ cuts of value at most α and $|V| = g$ by definition the lemma follows. $\qquad\square$

7.3 Upper Bound On CREP

We start the analysis of the lower bound of CREP by introducing several notions that we will use throughout the analysis.

We define the set $\text{DEL}(\sigma)$ as the set of components that were deleted by CREP during its execution given the input sequence σ.

We define the following notions for a component $C \in \text{DEL}(\sigma)$, i.e. the subgraph induced by the nodes of C has connectivity at least α and C consists of more than k nodes:

Let $\text{EPOCH}(C)$ denote the (node, time) pairs of nodes in C starting at the time after the time $\tau(node)$ when *node* was last turned into a singleton component, i.e.

$$\text{EPOCH}(C) = \bigcup_{n \in nodes(C)} \{n\} \times \{\tau(n) + 1, ..., \tau(C)\}.$$

This definition of an epoch is very similar to the one from [5] and [2]. Note that for $C \in \text{DEL}(\sigma)$, $\tau(C)$ denotes both the time of the creation as well as the time of deletion of C. We can use this definition of a component epoch $\text{EPOCH}(C)$ to uniquely assign each node to a deleted component C at each point in time t (except for nodes in components that persist until the end of sequence σ).

We assign all requests to $\text{EPOCH}(C)$ whose corresponding requests are deleted because of the deletion of component C and call the set of those requests $\text{REQ}(C)$. We split the requests from $\text{REQ}(C)$ into two sets: $\text{CORE}(C)$ contains all requests for which both nodes have already been assigned to C at the time of the request, i.e.

$$\text{CORE}(C) = \{r = \{u, v\} \in \sigma \,|\, (u, \text{TIME}(r)) \in \text{EPOCH}(C) \text{ and } (v, \text{TIME}(r)) \in \text{EPOCH}(C)\}.$$

These are the requests that led to the creation of component C. $\text{HALO}(C)$ contains all requests from $\text{REQ}(C)$ for which exactly one end point was associated with C at the time of the request. Note that this means that $\text{HALO}(C) = \text{REQ}(C) \backslash \text{CORE}(C)$.

These definitions allow us to differentiate between the highly-connected sub-graph induced by the nodes of C which are connected by requests from $\text{CORE}(C)$ and the edges leaving C from $\text{HALO}(C)$ which are relatively less dense as CREP has not merged any outer node with the component.

We start the analysis by bounding the communication cost of CREP that is due to serving requests from $\text{CORE}(C)$ for $C \in \text{DEL}(\sigma)$.

Lemma 8. With augmentation $2+\epsilon$, CREP pays at most communication cost $|C| \cdot \alpha$ for requests in $\text{CORE}(C)$ where $C \in \text{DEL}(\sigma)$.

Proof. First note that due to Lemma 4 CREP merges mergeable component sets as soon as they emerge. Whenever CREP performs a merge of a mergeable component set S, Lemma 7 states that there was at most total edge weight $(|S|-1) \cdot \alpha$ between the merged components, i.e. $w(S) \leq (|S|-1) \cdot \alpha$. Each such merge decreases the number of components that need to be merged in order to form component C by $|S|-1$. Hence CREP has payed at most $|C| \cdot \alpha$ communication cost for requests in $\text{CORE}(C)$. $\qquad\square$

We define $\text{FIN-WEIGHTS}(\sigma)$ as the total amount of edge weight between the components $\text{FIN-COMPS}(\sigma)$ which are present after the execution of CREP given input sequence σ.

Together with the fact that CREP pays for all requests in $\text{HALO}(C)$ for deleted components C we use these definitions as well as the previous lemma to bound the total communication cost of CREP in the following lemma.

Lemma 9. The cost of serving communication requests that CREP has to pay, denoted by $\text{CREP}^{req}(\sigma)$ given input sequence σ is bounded by

$$\text{CREP}^{req}(\sigma) \leq \sum_{C \in \text{DEL}(\sigma)} (|C| \cdot \alpha + |\text{HALO}(C)|) + \sum_{C \in \text{FIN-COMPS}(\sigma)} |C| \cdot \alpha + \text{FIN-WEIGHTS}(\sigma).$$

Proof. The number of communication requests that led to the creation of a component C is bounded by $|C| \cdot \alpha$ due to Lemma 7. If component C was deleted by CREP then also the edge weights corresponding to requests from $\text{HALO}(C)$ were reset to zero. All other edge weights were not changed. The remaining communication requests that have not been accounted for so far have either led to the creation of component $C \in \text{FIN-COMPS}(\sigma)$ and are hence also bounded by $|C| \cdot \alpha$ or have not led CREP to any merge and are hence contained in $\text{FIN-WEIGHTS}(\sigma)$. This concludes the proof. $\qquad\square$

We continue our analysis by bounding the migration cost of CREP in the following lemma.

Lemma 10. With augmentation $2 + \epsilon$, CREP pays at most migration costs of

$$\text{CREP}^{mig}(\sigma) \leq \sum_{C \in \text{DEL}(\sigma) \cup \text{FIN-COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha.$$

Proof. First note that CREP only performs migrations when it merges components. We fix a component $C \in \text{DEL}(\sigma) \cup \text{FIN-COMPS}(\sigma)$ and bound the number of times each node of C is moved as CREP processes the requests that led to the creation of C.

As CREP only reserves additional space $\lfloor \epsilon \cdot |B| \rfloor$ for components of size at least $2/\epsilon$ for each component B and only moves component B when a merge results in a component of size more than $(1+\epsilon) * |B|$ each node of C is moved at most $(2/\epsilon + 1) + \log k$ times. Summing over all nodes in C that were actually moved by CREP bounds the number of migrations by $|C| \cdot ((2/\epsilon + 1) \log k)$ as components get deleted without migrations once they contain more than k nodes. This leads to the desired bound on the migration costs as each node migration incurs cost α to CREP. $\qquad\square$

Finally we summarize our results from Lemma 9 and Lemma 10 in the following lemma inorder to obtain the final upper bound on the cost of CREP.

Lemma 11. With augmentation $2 + \epsilon$, CREP pays at most total cost

$$2 \cdot \sum_{C \in \text{COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha + \sum_{C \in \text{DEL}(\sigma)} |\text{HALO}(C)| + \text{FIN-WEIGHTS}(\sigma).$$

where $\text{COMPS}(\sigma) = \text{DEL}(\sigma) \cup \text{FIN-COMPS}(\sigma)$.

Proof. We sum the results from Lemma 9 and Lemma 10 to obtain the lemma:

$$
\begin{aligned}
\text{CREP}(\sigma) &\leq \text{CREP}^{req} + \text{CREP}^{mig} \\
&\leq \sum_{C \in \text{DEL}(\sigma)} (|C| \cdot \alpha + |\text{HALO}(C)|) + \sum_{C \in \text{FIN-COMPS}(\sigma)} |C| \cdot \alpha + \text{FIN-WEIGHTS}(\sigma) \\
&\quad + \sum_{C \in \text{COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha \\
&\leq 2 \cdot \sum_{C \in \text{COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha + \sum_{C \in \text{DEL}(\sigma)} |\text{HALO}(C)| \\
&\quad + \text{FIN-WEIGHTS}(\sigma).
\end{aligned}
$$

$\qquad\square$

7.4 Lower Bound on OPT

In this section we bound the cost on OPT by assigning cost to OPT based on the size of the components C that CREP deletes and the associated adjacent edges $\text{HALO}(C)$ which CREP resets to zero during the deletion of C. In order to achieve this we first introduce some additional notions.

First we define the term *offline interval* of a node v to be the time between two migrations of v in the solution of OPT. More specifically an offline interval of node v either starts at time zero (if it is the first offline interval of v) or after a migration of v and ends with the next migration of node v that OPT performs.

Furthermore we say that an offline interval is contained in the epoch $\text{EPOCH}(C)$ of a component $C \in \text{DEL}(\sigma)$ if it ends before the time $\tau(C)$. Note that $\tau(C)$ is both

the time of the creation of C in the solution of CREP and the time of its deletion as $C \in \text{DEL}(\sigma)$.

We assign a request r involving the node v to an offline interval of v if it is both the first offline interval of one of the end points of r that ends and if the offline interval ends before the deletion of the edge representing r due to a component deletion.

The requests from $\mathcal{H} = \bigcup_{C \in \text{DEL}(\sigma)} \text{HALO}(C)$ that are not assigned to any offline interval are then those which are deleted due to the deletion of a component that took place before the corresponding offline interval ended.

We start by bounding the total edge weight (the total number of requests) we assign to any one offline interval when limiting ourselves to requests from $\mathcal{H}$ which CREP pays for but OPT does not. We denote the requests in question by N, i.e. those from $N = \mathcal{H} \backslash P$. Note that $\mathcal{H}$ only contains requests which CREP payed for due to the definition of $\text{HALO}(C)$.

Lemma 12. We assign at most $k \cdot \alpha$ requests from N to any one offline interval.

Proof. We fix an arbitrary offline interval of node v. Observe that none of the nodes involved in the assigned requests are moved by OPT during the offline interval, hence all the requests in question involve only nodes that OPT has placed on the same server as v during the offline interval.

The number of such nodes is hence limited by the server capacity k. As we only examine requests from $\mathcal{H}$ we know that none of these requests have led CREP to perform any merges, hence there were at most α requests between v and any one of the other nodes on its server. This bounds the number of requests assigned to the offline interval by $k \cdot \alpha$. $\qquad\square$

In the following lemma we combine this result with a bound on the cost of OPT due to requests we have not accounted for, namely those from $\text{CORE}(C)$ and those that are not contained in an offline interval because they get deleted by CREP due to a component deletion that takes place before their offline interval ends.

Lemma 13. The cost of the solution of OPT given input sequence σ is bounded by

$$\text{OPT}(\sigma) \geq 1/2 \cdot \sum_{C \in \text{DEL}(\sigma)} |C|/k \cdot \alpha + |\text{HALO}(C)|/k.$$

Proof. Let P denote the set of edges from $\bigcup_{C \in \text{DEL}(\sigma)} \text{HALO}(C)$ that both CREP and OPT pay for and let I denote the set of requests we have assigned to offline intervals.

For the following part of the proof we fix an arbitrary component $C \in \text{DEL}(\sigma)$. Let R denote the set of requests from $\text{HALO}(C)$ that were not assigned to any offline interval. This means that the nodes involved in requests from R were not moved during the processing of requests from R until the time of deletion of C.

The number of nodes contained in C or connected to C via edges representing requests from R is at least $|C| + R/\alpha$ since requests from R have not led CREP to perform any migrations. Because of this fact OPT must have placed those nodes on

14

at least $\frac{|C|+R/\alpha}{k}$ different servers. As OPT does not pay for any requests from R it follows that OPT must have placed the nodes from C in $\frac{|C|+R/\alpha}{k}$ different servers.

We first examine the case in which OPT does not move any nodes from C during EPOCH(C). In this case OPT must partition a graph containing the nodes from C which are connected via edges representing the requests from CORE(C)into migrations. Because of this fact OPT must have placed those nodes on at least $\frac{|C|+R/\alpha}{k}$ parts. As CREP merged component C this graph is α-connected and hence Lemma 6 gives that OPT has to cut at least edges of total weight migrations. Because of this fact OPT must have placed those nodes on at least $\frac{|C|+R/\alpha}{k} \cdot \alpha = |C|/k \cdot \alpha + R/k$.

For the more general case in which OPT may perform node migrations during EPOCH(C) we adapt the graph construction from above as follows: we add a vertex representing each (node, time) pair from EPOCH(C). We connect each (node, time) pair p with edges of weight α to the pairs of the same node that represent the time step directly before and directly after p (if they exist in the graph). These edges represent the fact that OPT may choose to migrate a node between any two time steps in EPOCH(C). Additionally we add an edge of weight one for each request $r = \{u, v\}$ from CORE(C) by connecting the nodes in the graph that represent the pairs (u, t) and (v, t), respectively. OPT once again has to partition this graph into $\frac{|C|+R/\alpha}{k}$ parts.

Note that we only added edges of weight α to the graph and hence this graph is also α-connected. We conclude that once again OPT has to cut edges of weight at least $\frac{|C|+R/\alpha}{k} \cdot \alpha = |C|/k \cdot \alpha + R/k$.

Finally we need to account for the fact that the migrations of nodes from C that OPT performs also end offline intervals and might hence be accounted for twice in our analysis up to this point:

$$2 \cdot \text{OPT}(\sigma) \geq \sum_{(C \in \text{DEL}(\sigma))} |C|/k \cdot \alpha + R(C)/k + I/(k \cdot \alpha) \cdot \alpha + |P| + |OP|$$

$$\geq \sum_{(C \in \text{DEL}(\sigma))} |C|/k \cdot \alpha + \text{HALO}(C)/k$$

where the last equality follows from the fact that $\bigcup_{C \in \text{DEL}(\sigma)} R(C) \cup I \cup P \cup OP = \bigcup_{C \in \text{DEL}(\sigma)} \text{HALO}(C)$ as the different $R(C)$ as well as I, P and OP are disjoint. Hence the cost of OPT is at least $1/2 \cdot \sum_{(C \in \text{DEL}(\sigma))} |C|/k \cdot \alpha + \text{HALO}(C)/k$ as OPT pays for requests from P by the definition of P. $\qquad\square$

7.5 Competitive Ratio

In this section we combine the results of Lemma 11 and Lemma 13 to obtain the following theorem which gives us the desired competitive ratio.

Theorem 14. With augmentation $(2 + \epsilon)$ the competitive ratio of CREP is in $O(k \log k)$.

Proof. We arbitrarily fix an input sequence σ and use our previous results to bound the competitive ratio of CREP. We define COMPS(σ) := DEL(σ)∪FIN-COMPS(σ) in

order to improve readability. Let P denote the set of edges from $\bigcup_{C \in \mathrm{DEL}(\sigma)} \mathrm{HALO}(C)$ that both CREP and OPT pay for.

$$\frac{\mathrm{CREP}(\sigma) - \mathrm{FIN\text{-}WEIGHTS}(\sigma)}{\mathrm{OPT}(\sigma)}$$

$$\leq \frac{2 \cdot \sum_{C \in \mathrm{COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha + \sum_{C \in \mathrm{DEL}(\sigma)} |\mathrm{HALO}(C)|}{1/2 \cdot \sum_{C \in \mathrm{DEL}(\sigma)} |C|/k \cdot \alpha + |\mathrm{HALO}(C)|/k + |P|}$$

$$\leq k \log k \frac{2 \cdot \sum_{C \in \mathrm{DEL}(\sigma)} |C| \cdot (2/\epsilon + 1) \cdot \alpha + \sum_{C \in \mathrm{DEL}(\sigma)} |\mathrm{HALO}(C)|}{1/2 \sum_{(C \in \mathrm{DEL}(\sigma)} |C| \cdot \alpha/2 + |\mathrm{HALO}(C)|)} + \beta$$

$$= O(k \log k) + \beta$$

where

$$\beta = \sum_{C \in \mathrm{FIN\text{-}COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha$$

Let $\beta' = \beta + \mathrm{FIN\text{-}WEIGHTS}(\sigma)$. Then it follows that

$$\frac{\mathrm{CREP}(\sigma)}{\mathrm{OPT}(\sigma)} \leq O(k \log k) + \beta'.$$

To obtain the bound on β we observe that the components in FIN-COMPS(σ) each are of size at most k since they were not deleted by CREP. This allows us to derive to bound $\sum_{C \in \mathrm{FIN\text{-}COMPS}(\sigma)} |C| \cdot ((2/\epsilon + 1) + \log k) \leq l \cdot k \cdot ((2/\epsilon + 1) + \log k)$. Since at the end of the execution of CREP there can be at most $k \cdot l$ components, Lemma 7 allows us to bound FIN-WEIGHTS(σ) by $k \cdot l \cdot \alpha$. Hence we conclude that $\beta \leq l \cdot k \cdot ((2/\epsilon + 1) + \log k) \cdot \alpha + k \cdot l \cdot \alpha \in O(k \log k)$. $\qquad \square$

8 Implementation Details

In this section we first describe the ideas behind our implementation of CREP and then illustrate them by discussing pseudocode of the implemented algorithm.

8.1 Algorithm Explanations

In this section we describe our implementation of CREP with augmentation $2 + \epsilon$ in greater detail.

In order to limit the section of the graph G maintained by CREP that needs to be updated upon a new request between nodes of different components we maintain a decomposition tree defined as follows: the root represents the whole graph and is assigned the connectivity of the whole graph. Given a node v in the tree that represents a (sub-)graph G' of G, we decompose G' into sub-graphs whose connectivity is strictly larger than that of G' and add children to v for each such sub-graph. We

Figure 3: example graph

do not decompose sub-graphs of connectivity at least α any further as we only need to identify whether a new sub-graph of connectivity at least α was created by the insertion of the most recent request.

Figure 4 illustrates this decomposition for the graph shown in Figure 3. In the decomposition tree we have labelled each node with the corresponding subset of vertices and the connectivity of the graph induced by these vertices.

If a new request is revealed to CREP then we only need to update the smallest subtree of the decomposition tree which still contains both end points of the request. For example in the case of a new request between 0 and 1 we only need to recompute the decomposition of the sub-graph induced by the vertex set $\{0, 1, 2\}$.

This is correct because we can view each decomposition of a sub-graph G' into smaller graphs of a higher connectivity as a set of cuts that separates the nodes of G'. Inserting a new edge within G' may only increase the value of the cuts which result in the decomposition of G', but do not affect cuts separating G' from other sub-graphs.

If such a new request led to the creation of a new component this means that two old components that were at least α-connected were merged and hence the number of leaves in the decomposition tree decreased. If this is the case then the algorithm checks whether the new component contains more than k nodes. If this is the case then the component is deleted and split into singleton components, each containing one node from the deleted component.

Upon such a component deletion the edges inside of and adjacent to the component are deleted, i.e. their weight is reset to zero. This means that the decomposition tree needs to be recomputed in order to reflect this change.

If however the resulting component C contains at most k nodes the algorithm tries to collocate the nodes of the component while minimizing migration costs, i.e. looking for a cluster which contains as many nodes of the newly merged component as possible but which also has enough free capacity for the remaining nodes to be moved there and for additional reservation $\min\{k, k - \epsilon|C|\}$ if C contains at least $2/\epsilon$ nodes.

In the next section we provide more detailed pseudocode describing our implementation.

Figure 4: decomposition tree for the graph from Figure 3 for $\alpha = 4$

8.2 Algorithm Pseudocode

Algorithm 4 insertAndUpdate(a,b)

---
  if comp[a] == comp[b] **then**
    return
  end if
  addEdge(a, b)
  updateDecomposition(a, b)
  del← updateMapping$(alphaConnectedComponents)$
  delComponents(del)

Algorithm 4 is the main function that is called upon each new request. It checks whether the new request is between different α-connected components. If this is not the case it determines that this request can not change the decomposition and returns.

Otherwise the weight of the corresponding edge is increased and other routines are called that update the decomposition based on this new edge.

Algorithm 5 first determines the smallest sub-graph in the decomposition tree that contains both end points of the request and decomposes this sub-graph. For this decomposition step we use the algorithm proposed by Chang et al. ([6]). We explain this algorithm in greater detail in Section 8.3.

Afterwards the routine updateMapping is called which compares the number of components to the number of components before the arrival of the request. Only if the number of components has decreased it checks for the new component as otherwise there was no component merge.

If there was a merge then the routine checks is size and decides whether to delete or to collocate based on the logic described in the previous section.

If a deletion has to be performed then this step is done in the routine delComponents (Algorithm 6) which resets all edge weights both of edges between nodes of the component as well as all adjacent edges and finally starts the decomposition of

Algorithm 5 updateDecomposition(a,b)

---

q← findSmallestSubgraph(a, b)
while q not empty **do**
   current← q.popFront()
   if res.connectivity==alpha **then**
     continue
   end if
   res← decompose(current, current.connectivity+1)//decomposition based on s-t-cuts
   current.connectivity ← value of smallest encountered cut
   if current.connectivity≥alpha **then**
     continue
   end if
   childrenQueue ← res
   //make sure that only subgraphs with higher connectivity are added as children

   while childrenQueue not empty **do**
     c←childrenQueue.pop()
     cRes←decompose(c, current.connectivity+1)
     c.connectivity ← value of smallest encountered cut
     if decompose returned only one graph **then**
       current.children.add(cRes)
       if cRes has connectivity smaller than alpha **then**
         q.push(cRes)
       end if
     else
       childrenQueue.add(cRes)
     end if
   end while
end while

---

the whole graph in order to arrive at a new decomposition that follows the definition from the previous section.

Algorithm 6 delComponents(del)

---

    delAllEdges(del)
    root.connectivity=0
    root.children={}
    updateDecomposition(0,1)

8.3   On the Decomposition of a Subgraph

In this section we describe our algorithm for the decomposition of a given subgraph represented by a node in the decomposition tree. Specifically we describe the general idea behind our implementation of the subroutine *decompose(treeNode, c)* which we have based on the algorithm described by Chang et al. ([6]). In this algorithm c denotes the connectivity that is to be the basisi of this decomposition step.

Given a node v of the decomposition tree, first a *partition graph* is constructed which is a graph consisting of the nodes in the subgraph represented by v and the edges which are between the nodes of the subgraph. This partition graph also supports merges and cuts of its nodes. More specifically the partition graph is initialized as a graph $P = (V, E)$ with $V = nodes(v)$ and $E = \{e = \{u, w\} \in E' | u \in nodes(v) \text{ and } w \in nodes(v)\}$ where E' represents the set of edges in the graph maintained by CREP. Additionally we maintain a mapping M which assigns each node from V a set of the nodes in the subgraph represented by the decomposition tree node v. Initially M assigns each node in V the subset containing only the node itself.

We now run a *maximum adjacency search* algorithm, sometimes also called *maximum cardinality search* algorithm ([18]) in order to obtain a minimum s-t-cut of the graph.

The maximum adjacency search algorithm is defined as follows: We start with an empty list L to which we add an arbitrary node of P. We then continually add the most tightly connected node from V to L, i.e. the node which is connected to the nodes in L via edges of the most total weight. Stoer and Wagner ([18]) have shown that the edges between the last two nodes s and t added to L form a minimum s-t-cut.

We use the value of this cut in order to decide whether to merge the nodes s and t or whether to separate them: if the cut has value less than c we separate the nodes, otherwise we perform a merge.

Here the separation of the nodes s and t means that we remove all edges in the cut from the edges E of the partition graph P.

In the case of a merge we combine the nodes s and t and merge the outgoing edges, i.e. we replace the set of nodes V of P by the set $V' = V \setminus \{s, t\} \cup \{v'\}$. The edges E of P are modified by removing all edges adjacent to s and t and adding an edge $e' = \{v', u\}$ of weight $w(\{s, u\}) + w(\{t, u\})$ where $w(e)$ denotes the weight of

edge e if it exists and is equal to zero otherwise. Furthermore we adjust the mapping M by setting $M(v') = M(s) \cup M(t)$.

We continually run this algorithm until P contains no edges, i.e. until $E = \emptyset$. The sets of nodes mapped to each ot the nodes of P by M now represent candidate subgraphs for the decomposition. Note though that we have only cut according to minimum s-t-cuts and not according to minimum cuts. This means that the specific sequence in which we have performed the cuts may influence the result. This can be remedied by repeating the procedure on the resulting subgraphs until it returns a subgraph of only one node. This is due to the fact that Chang et al. ([6]) have shown that this procedure always cuts a subgraph of connectivity less than c at least once (see Cutability Property in [6]).

In order to speed up this computation we use the heap data structure proposed and analyzed by Chang et al. ([6]).

Thus we conclude that this procedure correctly decomposes a given subgraph as Chang et al. have also stated in Theorem 3.1 in [6].

8.4 Time Complexity

In this section we discuss the time complexity of an update of the decomposition tree after a new request has been received. The worst case is in this case given when the whole tree has to be recomputed.

We first discuss the time complexity of decomposing a single tree node v. Let the corresponding subgraph be denoted by $G_v = (V_v, E_v)$. Since each iteration of the subroutine *decompose* performs at least one cut as long as the connectivity of the given graph is smaller than the current threshold c we conclude that after at most $|V_v|$ iterations of decompose a correct decomposition is found. Each step of *decompose* can be performed in $O(|V_v| \cdot |E_v|)$ a maximum adjacency search finds an arbitrary minimum s-t-cut in time $O(|E_v|)$ as shown in theorem 4.1 in [6] and as there are at most $|V_v|$ minimum s-t-cuts computed for each invocation of *decompose*.

Hence the complexity of decomposing the subgraph represented by a tree node v is in $O(|V_v|^2|E_v|)$. Let C_v denote the time needed for the decomposition of the subgraph represented by decomposition tree node v.

We now sum this complexity over the nodes for each connectivity level of the decomposition tree. To this end let level(i) denote all nodes in the decomposition tree which are of connectivity exactly i.

$$\sum_{i=0}^{\alpha} \sum_{v \in \text{level}(i)} C_v \leq \sum_{i=0}^{\alpha} O(|V|^2|E|) \in O(\alpha|V|^2|E|).$$

We conclude our analysis of the time complexity by observing the polynomial-time complexity of $O(\alpha|V|^2|E|)$ we have just shown.

9 Evaluation

In this section we evaluate the quality of the results and the performance of our algorithm implementation by comparing it with several algorithms described in Sec-

on input data sets described in .

9.1 Algorithms

We first present the algorithms we compare in order to evaluate our results.

We consider both algorithmic approaches based on maintaining a second-order partition of the nodes into components we discussed in Section 6, i.e. ADJ-DEL where on a component deletion all edges inside of and adjacent of the component are deleted and CORE-DEL which only deletes internal edges.

We compare these algorithms both with the static graph partitioning algorithm METIS_PartGraphRecursive implemented in the METIS framework ([13, 14]) which we will refer to from now on as STATIC and the adaptive/dynamic algorithm ParMETIS_V3_AdaptiveRepart ([15, 17, 16]) available in the ParMetis framework, referred to as ADAPT. Both frameworks are known to produce very good results and to be very fast.

9.2 Input Data

As input data we use several HPC traces (Mocfe, NeckBone, MultiGrid), the nature of the data is described in more detail by Avin, Ghobadi, Griner and Schmid in [4]. For the sake of readability we use abbreviations NB for NeckBone and MG for MultiGrid.

All data sets contain 1024 different communication nodes and are limited to the first 300 000 requests. The value of α is set to be 6 and the algorithm was tasked to partition the nodes into 32 clusters of size 32 each, i.e. $k = l = 32$. The dynamic algorithms are allowed to use augmentation with a factor of 2.1, i.e. for the dynamic algorithms the maximum cluster capacities are $\lfloor 32 \cdot 2.1 \rfloor = 67$.

9.3 Comparison of ADJ-DEL and CORE-DEL

We compare the results of ADJ-DEL and CORE-DEL first. Figure 5 shows the resulting total cost of both algorithms on the different input sets as well as how this cost is split between communication and migration costs. One can see that ADJ-DEL always produces better results, especially for the second and third input set.

We now discuss the differences of the two algorithms in terms of communication and migration cost. The figures illustrate that ADJ-DEL pays only very slightly less for communication requests than CORE-DEL whereas the former pays significantly reduced migration cost. This suggests that that the deletion of adjacent edges indeed improves the quality of the results by reducing the number of migrations.

Finally Figure 6 compares the running times of both algorithms. One can see that ADJ-DEL performs drastically better than CORE-DEL in this regard. This may be due to the fact that ADJ-DEL deletes edges more frequently and thus needs to take less edges into account when updating the decomposition tree.

Also note that both algorithms pay significantly more for migrations than for communication. This suggests that there may be room for fine-tuning these algorithms in order achieve a more balanced distribution of the cost. One such adjustment

Figure 5: comparison of the total cost of ADJ-DEL (left) and CORE-DEL (right)

is investigated in the next section, namely whether we can improve the results by changing the algorithms in such a way that they only merge once a connectivity of $2 \cdot \alpha$ is reached.

9.4 On the Influence of α

In the previous section we observed that both CORE-DEL and ADJ-DEL have higher migration cost than communication cost. This leads to the questions whether we can find adjustments that allow us to achieve a better balance of the different costs. In this section we investigate the influence of α on the quality of the solution. Namely we discuss the results of both algorithms for the case where they only merge components once a connectivity of at least $2 \cdot \alpha$ is reached. Note that this only affects the analysis by constant factors, namely the statements from Lemma 6 and Lemma 7 are multiplied by a factor of two.

Figure 7 shows the total cost of ADJ-DEL and CORE-DEL for this scenario. One can see that the cost of both algorithms is reduced, but CORE-DEL has improved significantly more than ADJ-DEL.

When looking at the distribution of this cost to communication and migration costs one can see that now the communication costs of both algorithms is higher than the migration cost in all cases. For CORE-DEL this has improved the balance of both costs, but ADJ-DEL now has significantly higher communication cost than migration cost. This suggests that by adjusting the exact value of connectivity at which these algorithms perform their merges one might be able to further improve the quality of the results. It may even be possible to adjust this and parameter dynamically based on some data gathered as requests are processed.

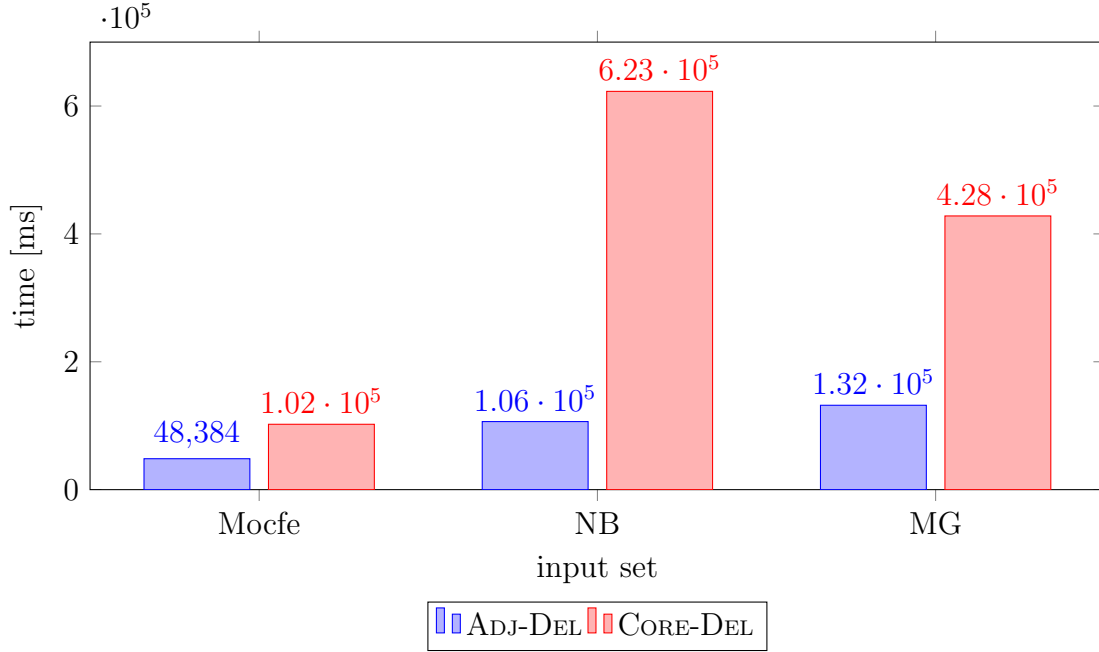These improvements come at a cost of running time as Figure 8 shows, almost

Figure 6: comparison of running time of Adj-Del and Core-Del

doubling the times of the previous results in some cases. Especially the already worse run times of Core-Del are drastically increased.

9.5 Results of Adapt and Static

Figure 9 shows the total cost of the solutions of Adapt and Static. One can see that Static is able to produce results that far surpass Adapt. In fact the results of Static are also a significant improvement over the results presented before. But it is also important to note that Static is only useful in scenarios where the communication patterns and frequencies stay mostly the same over time. Otherwise the approach to record the communication graph and then use Static in order to compute a partitioning may lead to consistent results. We stress that the implementations of Static and Adapt rely on heuristics only and do not provide any guarantees while we have shown that our algorithm Adj-Del is a competitive algorithm with competitive ratio $O(k \log k)$.

Finally Figure 10 illustrates the run times of Adapt and Static. Note that static is only run once on the communication graph after all requests have been revealed which means that it is expected for Static to have the fastest running time. But also Adapt achieves faster running times than our dynamic implementations.

10 Conclusion

We have discussed different ideas for component-based algorithms for the Dynamic Balanced Graph Partitioning Problem.

We have shown that our algorithm Adj-Del has a competitive ratio of $O(k \log k)$ and we have presented our decomposition based approach to its implementation.

Figure 7: comparison of the total cost of Adj-Del (left) and Core-Del (right) in the case where merges are performed at connectivity $2 \cdot \alpha$


Figure 8: comparison of running time of Adj-Del and Core-Del in the case where merges are performed at connectivity $2 \cdot \alpha$
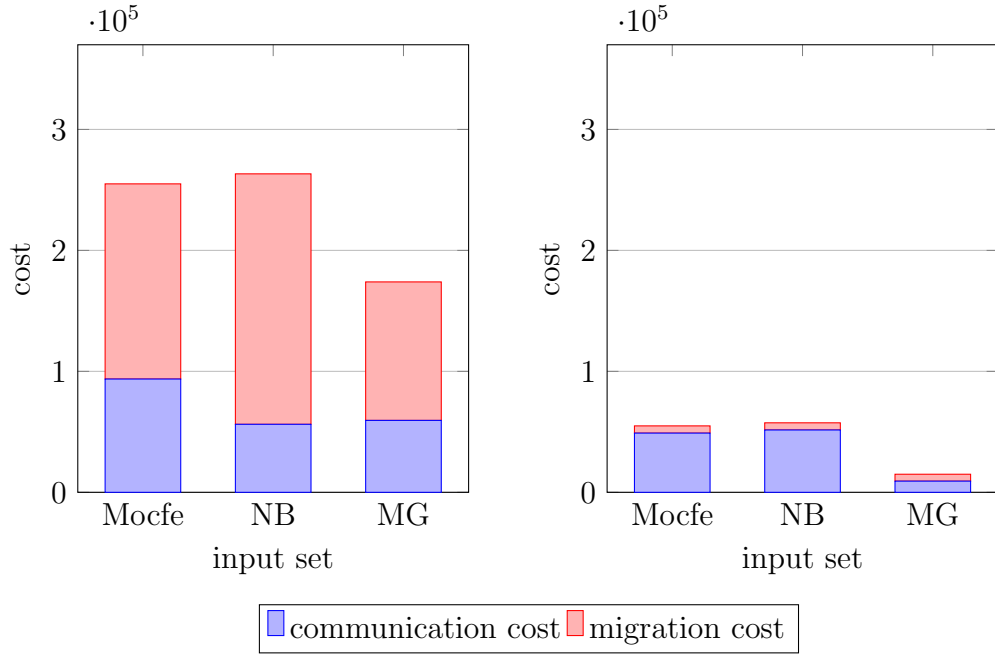
Figure 9: comparison of the total cost of ADAPT (left) and STATIC (right)

Figure 10: comparison of running time of ADAPT and STATIC

Finally we have evaluated the results of our implementation by comparing it with three other algorithms: Core-Del which is a variation of Adj-Del which deletes fewer edges and Adapt and Static which are available via METIS and ParMETIS frameworks, respectively. Furthermore we have explored the possibility of making slight adaptations of our algorithms in order to improve their results.

11   Future Work

Even though we have already shown very promising results as to the overall performance of our algorithms together with theoretical guarantees there are still areas in which our work could be expanded upon.

As discussed in Section 9.4 we see great potential in making adjustments to our algorithms, especially Core-Del, in order to improve their results even further while also preserving our result on the competitive ratio.

There may also be potential in adapting and improving our decomposition tree data structure in order to improve running times.

One such avenue for improvement may be to parallelize the computation of the decomposition of disjoint subgraphs. This could be both used to handle multiple requests in parallel if they only affect disjoint sub-graphs, i.e. their corresponding trees do not overlap in the decomposition tree, as well as in order to speed up the computation of the decomposition the children trees during the update process after a new request has been received by Crep.

References

[1] Konstantin Andreev and Harald Räcke. Balanced Graph Partitioning. *Theory of Computing Systems*, 39(6):929–939, oct 2006.

[2] Chen Avin, Marcin Bienkowski, Andreas Loukas, Maciej Pacut, and Stefan Schmid. Dynamic Balanced Graph Partitioning. *no idea*, 2015.

[3] Chen Avin, Louis Cohen, Mahmoud Parham, and Stefan Schmid. Competitive clustering of stochastic communication patterns on a ring. *Computing*, 101(9):1369–1390, sep 2018.

[4] Chen Avin, Manya Ghobadi, Chen Griner, and Stefan Schmid. Measuring the Complexity of Packet Traces. *arXiv:1905.08339v1*, 2019.

[5] Chen Avin, Andreas Loukas, Maciej Pacut, and Stefan Schmid. Online Balanced Repartitioning. In *Lecture Notes in Computer Science*, pages 243–256. Springer Berlin Heidelberg, 2016.

[6] Lijun Chang, Jeffrey Xu Yu, Lu Qin, Xuemin Lin, Chengfei Liu, and Weifa Liang. Efficiently computing k-edge connected components via graph decomposition. In *Proceedings of the 2013 international conference on Management of data - SIGMOD '13*. ACM Press, 2013.

[7] Leah Epstein, Csanád Imreh, Asaf Levin, and Judit Nagy-Gyrgy. On variants of file caching. In *Automata, Languages and Programming*, pages 195–206. Springer Berlin Heidelberg, 2011.

[8] Amos Fiat, Richard Karp, Mike Luby, Lyle McGeoch, Daniel Sleator, and Neal E. Young. Competitive Paging Algorithms. *arXiv preprint cs/0205038*.

[9] Monia Ghobadi, Daniel Kilper, Ratul Mahajan, Amar Phanishayee, Nikhil Devanur, Janardhan Kulkarni, Gireeja Ranade, Pierre-Alexandre Blanche, Houman Rastegarfar, and Madeleine Glick. ProjecToR: Agile Reconfigurable Data Center Interconnect. In *Proceedings of the 2016 conference on ACM SIGCOMM 2016 Conference - SIGCOMM '16*. ACM Press, 2016.

[10] Navid Hamedazimi, Zafar Qazi, Himanshu Gupta, Vyas Sekar, Samir R. Das, Jon P. Longtin, Himanshu Shah, and Ashish Tanwer. FireFly. In *Proceedings of the 2014 ACM conference on SIGCOMM - SIGCOMM '14*. ACM Press, 2014.

[11] Erez Hartuv and Ron Shamir. A clustering algorithm based on graph connectivity. *Information Processing Letters*, 76(4-6):175–181, dec 2000.

[12] Monika Henzinger, Stefan Neumann, and Stefan Schmid. Efficient Distributed Workload (Re-)Embedding. *no idea*, 2019.

[13] George Karypis and Vipin Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM Journal on Scientific Computing*, 20(1):359–392, jan 1998.

[14] George Karypis and Vipin Kumar. Multilevelk-way partitioning scheme for irregular graphs. *Journal of Parallel and Distributed Computing*, 48(1):96–129, jan 1998.

[15] George Karypis and Vipin Kumar. Parallel multilevel series k-way partitioning scheme for irregular graphs. *SIAM Review*, 41(2):278–300, jan 1999.

[16] K. Schloegel, G. Karypis, and V. Kumar. A unified algorithm for load-balancing adaptive scientific simulations. In *ACM/IEEE SC 2000 Conference (SC'00)*. IEEE, 2000.

[17] Kirk Schloegel, George Karypis, and Vipin Kumar. Multilevel diffusion schemes for repartitioning of adaptive meshes. *Journal of Parallel and Distributed Computing*, 47(2):109–124, dec 1997.

[18] Mechthild Stoer and Frank Wagner. A simple min-cut algorithm. *Journal of the ACM*, 44(4):585–591, jul 1997.